

A3: Predicting Car Price (Classification)

GitHub repository: https://github.com/WinNawee/A3-Predicting_Car_Price

Live site: <http://ml.brain.cs.ait.ac.th:8031> (reachable from inside CSIM)

In A1 and A2 we predicted the exact selling price. In this assignment the same used car dataset is turned into a **4 class classification** problem:

`selling_price` is split into 4 price bands, and a multinomial Logistic Regression written from scratch predicts which band a car belongs to.

What this notebook covers:

- **Task 1:** put `selling_price` into 4 classes, extend the `LogisticRegression` class (softmax + gradient descent), and add `accuracy`, `precision`, `recall`, `f1_score`, macro and weighted averages to the class. I check my results against `sklearn.metrics.classification_report`.
- **Task 2:** add an optional Ridge (L2) penalty to the loss and gradient.
- **Task 3:** log the experiments with MLflow, register the best model in Staging, and set up CI/CD for the web app. The course MLflow server was down, so following the TA announcement (1 Oct 2026) the runs are logged to a local MLflow instance, like in A2. Screenshots of the MLflow UI are included in Sections 5 to 7, and the CI/CD and deployment are described in Section 8.

```
In [1]: import re
import time

import numpy as np
import pandas as pd
import matplotlib.pyplot as plt

from sklearn.model_selection import train_test_split
from sklearn.compose import ColumnTransformer
from sklearn.pipeline import Pipeline
from sklearn.impute import SimpleImputer
from sklearn.preprocessing import StandardScaler, OneHotEncoder
from sklearn.metrics import classification_report as
sk_classification_report

np.random.seed(42)
STUDENT_ID = "st127031"
```

1. Load the Car Price dataset

Same raw `Cars.csv` as A1 and A2. Columns: `name`, `year`, `selling_price`, `km_driven`, `fuel`, `seller_type`, `transmission`, `owner`, `mileage`, `engine`, `max_power`, `torque`, `seats`.

```
In [2]: df = pd.read_csv("data/Cars.csv")
print(df.shape)
df.head()
```

(8128, 13)

```
Out[2]:
```

	name	year	selling_price	km_driven	fuel	seller_type	transmission	owner	mileage	engine	max_power	torque	seats
0	Maruti Swift Dzire VDI	2014	450000	145500	Diesel	Individual	Manual	First Owner	23.4 kmpl	1248 CC	74 bhp	190Nm@ 2000rpm	5.0
1	Skoda Rapid 1.5 TDI Ambition	2014	370000	120000	Diesel	Individual	Manual	Second Owner	21.14 kmpl	1498 CC	103.52 bhp	250Nm@ 1500-2500rpm	5.0
2	Honda City 2017-2020 EXi	2006	158000	140000	Petrol	Individual	Manual	Third Owner	17.7 kmpl	1497 CC	78 bhp	12.7@ 2,700(kgm@ rpm)	5.0
3	Hyundai i20 Sportz Diesel	2010	225000	127000	Diesel	Individual	Manual	First Owner	23.0 kmpl	1396 CC	90 bhp	22.4 kgm at 1750-2750rpm	5.0
4	Maruti Swift VXi BSIII	2007	130000	120000	Petrol	Individual	Manual	First Owner	16.1 kmpl	1298 CC	88.2 bhp	11.5@ 4,500(kgm@ rpm)	5.0

```
In [3]: df.isna().sum()
```

```
Out[3]: name          0
year          0
selling_price  0
km_driven     0
fuel          0
seller_type   0
transmission  0
owner         0
mileage      221
engine        221
max_power     215
torque        222
seats         221
dtype: int64
```

2. Preprocessing (same steps as A1/A2, with one difference)

- `mileage` ("23.4 kmpl"), `engine` ("1248 CC") and `max_power` ("74 bhp"): remove the unit and keep the number.
- `torque` is dropped. Its format is very inconsistent and A1/A2 did not use it either.
- `brand` is the first word of `name`. Brands with fewer than 20 cars in the training set become "Other".
- Only Diesel and Petrol cars are kept. CNG/LPG are about 1% of the data and use a different mileage unit (km/kg), and they were removed in A1/A2 as well.
- `owner` is encoded as an ordered number (First Owner = 0, ..., Test Drive Car = 4).
- **Difference from A2:** A2 removed the Test Drive Car rows because their prices are unusually high. Here they are kept. There are only 5 of them (about 0.07% of the data), all priced in the Luxury band, so they have almost no effect on the results.
- Missing values are filled later inside the `ColumnTransformer` (median for numbers, most frequent value for categories), fitted on the training set only.

```
In [4]: def parse_number(s):
        """'23.4 kmpl' -> 23.4, '1248 CC' -> 1248.0, NaN stays NaN"""
        if pd.isna(s):
            return np.nan
        m = re.search(r"[\d.]+", str(s))
        return float(m.group()) if m else np.nan

df = df[df["fuel"].isin(["Diesel", "Petrol"])]

df["mileage"] = df["mileage"].apply(parse_number)
df["engine"] = df["engine"].apply(parse_number)
df["max_power"] = df["max_power"].apply(parse_number)
df["brand"] = df["name"].apply(lambda x: str(x).split(" ")[0])
df = df.drop(columns=["torque"])

OWNER_SCALE = {
    "First Owner": 0,
    "Second Owner": 1,
    "Third Owner": 2,
    "Fourth & Above Owner": 3,
    "Test Drive Car": 4,
}
df["owner"] = df["owner"].map(OWNER_SCALE)

print(df.shape)
df.isna().sum()
```

(8033, 13)

```
Out[4]: name          0
        year          0
        selling_price  0
        km_driven     0
        fuel          0
        seller_type    0
        transmission   0
        owner         0
        mileage       214
        engine        214
        max_power     208
        seats         214
        brand         0
        dtype: int64
```

Remove duplicate rows before splitting (data leakage fix 1)

The dataset has many repeated listings. If the same car is in both the training set and the test set, the model is tested on rows it has already seen, and the test score looks better than it really is. So duplicates are removed **before** the split. I compare all the columns the model can see (features and target), which also catches cars that only differ in the variant `name`.

```
In [5]: MODEL_VISIBLE_COLS = ["year", "km_driven", "mileage", "engine",
                             "max_power", "seats",
                             "owner", "fuel", "transmission", "seller_type",
                             "brand", "selling_price"]

n_before = len(df)
df = df.drop_duplicates(subset=MODEL_VISIBLE_COLS).reset_index(drop=True)
print(f"Dropped {n_before - len(df)} duplicate rows, {len(df)} unique
rows remain")
```

Dropped 1220 duplicate rows, 6813 unique rows remain

Train / validation / test split (data leakage fix 2)

- **train (64%)**: used to fit everything (brand grouping, class edges, imputers, scaler, model weights).
- **validation (16%)**: used to compare models and hyperparameters (no penalty vs ridge, the MLflow runs, choosing the best run).
- **test (20%)**: used **only once** at the end, on the model already chosen with the validation set. If I picked the best model by its test score, the test set would be part of model selection and the final score would be too optimistic.

```
In [6]: df_trainval, df_test = train_test_split(df, test_size=0.2,
random_state=42)
df_train, df_val = train_test_split(df_trainval, test_size=0.2,
random_state=42)
print("train:", df_train.shape, " val:", df_val.shape, " test:",
df_test.shape)

# check: no car in the test set should also be in the training set
overlap = df_test.merge(df_train, on=MODEL_VISIBLE_COLS, how="inner")
print("test rows that also appear in train:", len(overlap))
```

```
train: (4360, 13) val: (1090, 13) test: (1363, 13)
test rows that also appear in train: 0
```

Group brands into "Other". Only brands with at least 20 cars in the training set keep their own name. Every other brand becomes "Other", including brands that never appear in the training set (they may still show up in validation, test, or in the web app). The list of common brands is saved with the model so the app does exactly the same thing.

```
In [7]: brand_counts = df_train["brand"].value_counts()
COMMON_BRANDS = sorted(brand_counts[brand_counts >= 20].index)

def group_brand(b):
    return b if b in COMMON_BRANDS else "Other"

df_train = df_train.copy()
df_val = df_val.copy()
df_test = df_test.copy()
for _d in (df_train, df_val, df_test):
    _d["brand"] = _d["brand"].apply(group_brand)

NUM_COLS = ["year", "km_driven", "mileage", "engine", "max_power",
"seats", "owner"]
TWO_LEVEL_COLS = ["fuel", "transmission"]
MULTI_LEVEL_COLS = ["seller_type", "brand"]
FEATURE_COLS = NUM_COLS + TWO_LEVEL_COLS + MULTI_LEVEL_COLS
```

2.1 Put selling_price into 4 classes

I use `pd.qcut` (quantile cut) on the **training set only**, then apply the same bin edges to the validation and test sets with `pd.cut`. This gives 4 classes of about the same size.

I did not use an equal width `pd.cut`. A few very expensive cars stretch the price range so much that more than 90% of the cars would end up in the first bin.

```
In [8]: train_classes, bin_edges = pd.qcut(
        df_train["selling_price"], q=4, labels=[0, 1, 2, 3], retbins=True
    )
y_train = train_classes.astype(int).values

# open the outer edges so every price in val/test falls into some bin
bin_edges_open = bin_edges.copy()
bin_edges_open[0] = -np.inf
bin_edges_open[-1] = np.inf

def to_class(prices):
    return pd.cut(prices, bins=bin_edges_open, labels=[0, 1, 2, 3]).astype(int).values

y_val = to_class(df_val["selling_price"])
y_test = to_class(df_test["selling_price"])

CLASS_LABELS = {0: "Budget", 1: "Mid-range", 2: "Premium", 3: "Luxury"}
CLASS_PRICE_RANGES = {
    int(c): f"₹{bin_edges[c]:,.0f} – ₹{bin_edges[c+1]:,.0f}" for c in
    range(4)
}
for c in range(4):
    print(f"Class {c} ({CLASS_LABELS[c]}): {CLASS_PRICE_RANGES[c]}")

print("\nTrain class balance:", np.bincount(y_train))
print("Val class balance: ", np.bincount(y_val))
print("Test class balance: ", np.bincount(y_test))
```

```
Class 0 (Budget): ₹30,000 – ₹250,000
Class 1 (Mid-range): ₹250,000 – ₹405,000
Class 2 (Premium): ₹405,000 – ₹635,000
Class 3 (Luxury): ₹635,000 – ₹7,200,000
```

```
Train class balance: [1170 1011 1093 1086]
Val class balance:   [311 242 261 276]
Test class balance:  [360 289 357 357]
```

2.2 ColumnTransformer (same structure as the A1/A2 prep pipeline)

- Numeric columns: fill missing with the median, then standardize.
- Two level categories (fuel, transmission): fill with the most frequent value, then one hot encode with `drop="if_binary"`.
- Multi level categories (seller_type, brand): fill with the most frequent value, then one hot encode. `handle_unknown="ignore"` means a brand the model has never seen will not crash the web app.

It is fitted on the **training set only**.

```

In [9]: prep = ColumnTransformer(
    transformers=[
        ("num", Pipeline([
            ("fillna", SimpleImputer(strategy="median")),
            ("scale", StandardScaler()),
        ]), NUM_COLS),
        ("two", Pipeline([
            ("fillna", SimpleImputer(strategy="most_frequent")),
            ("encode", OneHotEncoder(drop="if_binary",
handle_unknown="ignore")),
        ]), TWO_LEVEL_COLS),
        ("multi", Pipeline([
            ("fillna", SimpleImputer(strategy="most_frequent")),
            ("encode", OneHotEncoder(handle_unknown="ignore")),
        ]), MULTI_LEVEL_COLS),
    ],
    remainder="drop",
)

def add_intercept(X):
    return np.concatenate([np.ones((X.shape[0], 1)), X], axis=1)

def to_dense(X):
    return X.toarray() if hasattr(X, "toarray") else X

# fit on train ONLY, validation and test are only transformed
X_train =
add_intercept(to_dense(prepare.fit_transform(df_train[FEATURE_COLS])))
X_val = add_intercept(to_dense(prepare.transform(df_val[FEATURE_COLS])))
X_test = add_intercept(to_dense(prepare.transform(df_test[FEATURE_COLS])))

k = 4 # number of classes
n = X_train.shape[1] # number of features, including the intercept
m = X_train.shape[0] # number of training samples

# Y must be one hot with shape (m, k) for softmax regression
Y_train_onehot = np.zeros((m, k))
for c in range(k):
    Y_train_onehot[Y_train == c, c] = 1

print("X_train:", X_train.shape, " X_val:", X_val.shape, " X_test:",
X_test.shape,
      " n features (incl. intercept):", n)

```

```

X_train: (4360, 30) X_val: (1090, 30) X_test: (1363, 30) n features
(incl. intercept): 30

```

3. Task 1: Multinomial Logistic Regression and metrics from scratch

3.1 The `LogisticRegression` class

The base is the class from `02 - Multinomial Logistic Regression.ipynb` (softmax output, cross entropy loss, batch / minibatch / stochastic gradient descent). My changes:

1. **Classification metrics are now methods of the class** (Task 1):
`accuracy`, `precision`, `recall`, `f1_score`, `macro_precision`, `macro_recall`, `macro_f1`, `weighted_precision`, `weighted_recall`, `weighted_f1`, plus `classification_report` to print them like sklearn. They are class methods, so they work on a trained model (`model.accuracy(y, yhat)`) and also without one (`LogisticRegression.accuracy(y, yhat)`), which I need for the mock data check.
2. **`use_penalty` and `l`**: optional Ridge (L2) penalty (Task 2).
3. **The gradient is divided by the batch size `m`**. The original class uses `grad = X.T @ error` without `1/m`. That was fine for the small iris example (about 100 rows), but our training set has 4,360 rows, so the gradient became very large and needed a tiny `alpha` to not diverge. After dividing by `m` the step size does not depend on the batch size, so the same `alpha` works for batch and minibatch.
4. **Softmax subtracts the row maximum** before `exp()` to avoid overflow. The probabilities do not change.

```

In [10]: class RidgePenalty:
    """L2 penalty on the weight matrix W (n_features, n_classes).
    Row 0 of W is the bias, which is not penalized."""

    def __init__(self, l=0.0):
        self.l = l

    def __call__(self, W):
        return self.l * np.sum(np.square(W[1:, :]))

    def derivation(self, W):
        # derivative of l * sum(W^2) is 2 * l * W
        grad = self.l * 2 * W
        grad[0, :] = 0.0
        return grad

class NoPenalty:
    """Used when use_penalty=False, adds nothing to the loss or
    gradient."""

    def __call__(self, W):
        return 0.0

    def derivation(self, W):
        return np.zeros_like(W)

class LogisticRegression:

    def __init__(self, k, n, method="batch", alpha=0.001, max_iter=5000,
        use_penalty=False, l=0.1, verbose=False,
        random_state=None):
        self.k, self.n, self.method = k, n, method
        self.alpha, self.max_iter, self.verbose = alpha, max_iter,
        verbose
        self.use_penalty, self.l = use_penalty, l
        self.random_state = random_state
        self.regularization = RidgePenalty(l) if use_penalty else
        NoPenalty()

    # -----
    # training
    # -----

    def fit(self, X, Y):
        rng = np.random.RandomState(self.random_state)
        self.W = rng.rand(self.n, self.k)
        self.losses = []
        m_total = X.shape[0]
        self.m_total = m_total # used to scale the ridge term (see
        gradient)

        if self.method == "batch":
            start = time.time()
            for i in range(self.max_iter):
                loss, grad = self.gradient(X, Y)
                self.losses.append(loss)
                self.W = self.W - self.alpha * grad
                if self.verbose and i % 1000 == 0:

```

```

        print(f"iter {i}: loss={loss:.4f}")
    if self.verbose:
        print(f"time taken: {time.time()-start:.2f}s")

    elif self.method == "minibatch":
        start = time.time()
        batch_size = max(1, int(0.3 * m_total))
        for i in range(self.max_iter):
            ix = rng.randint(0, m_total)
            batch_X, batch_Y = X[ix:ix + batch_size], Y[ix:ix +
batch_size]

            if batch_X.shape[0] == 0:
                continue
            loss, grad = self.gradient(batch_X, batch_Y)
            self.losses.append(loss)
            self.W = self.W - self.alpha * grad
            if self.verbose and i % 1000 == 0:
                print(f"iter {i}: loss={loss:.4f}")
        if self.verbose:
            print(f"time taken: {time.time()-start:.2f}s")

    elif self.method == "sto":
        start = time.time()
        used = []
        for i in range(self.max_iter):
            idx = rng.randint(m_total)
            while idx in used:
                idx = rng.randint(m_total)
            X_i, Y_i = X[idx, :].reshape(1, -1), Y[idx].reshape(1,
-1)

            loss, grad = self.gradient(X_i, Y_i)
            self.losses.append(loss)
            self.W = self.W - self.alpha * grad
            used.append(idx)
            if len(used) == m_total:
                used = []
            if self.verbose and i % 1000 == 0:
                print(f"iter {i}: loss={loss:.4f}")
        if self.verbose:
            print(f"time taken: {time.time()-start:.2f}s")
    else:
        raise ValueError('method must be one of "batch", "minibatch"
or "sto".')
    return self

def gradient(self, X, Y):
    m = X.shape[0] # size of this batch
    N = getattr(self, "m_total", m) # size of the whole training
set
    h = self.h_theta(X, self.W)
    # cross entropy is averaged over the batch. The ridge term is
divided
    # by the full training size N, so its strength does not change
with
    # the batch size (for "batch" N == m, for "sto" m == 1).
    ce_loss = -np.sum(Y * np.log(h + 1e-15)) / m
    penalty_loss = self.regularization(self.W) / N
    loss = ce_loss + penalty_loss

```

```

        error = h - Y
        grad = (X.T @ error) / m + self.regularization.derivation(self.W)
    / N
    return loss, grad

def softmax(self, z):
    z = z - np.max(z, axis=1, keepdims=True)
    e = np.exp(z)
    return e / np.sum(e, axis=1, keepdims=True)

def h_theta(self, X, W):
    return self.softmax(X @ W)

def predict(self, X):
    return np.argmax(self.h_theta(X, self.W), axis=1)

def predict_proba(self, X):
    return self.h_theta(X, self.W)

# -----
# Task 1: classification metrics, written from scratch
# (no sklearn.metrics used in here)
# -----
@staticmethod
def _confusion_counts(y_true, y_pred, c):
    """TP, FP, FN for class c (one class vs the rest)."""
    y_true, y_pred = np.asarray(y_true), np.asarray(y_pred)
    tp = int(np.sum((y_true == c) & (y_pred == c)))
    fp = int(np.sum((y_true != c) & (y_pred == c)))
    fn = int(np.sum((y_true == c) & (y_pred != c)))
    return tp, fp, fn

@staticmethod
def _classes(y_true, y_pred, classes):
    """Classes to report on. Like sklearn, use every label that
appears
in y_true OR y_pred."""
    if classes is not None:
        return classes
    return np.unique(np.concatenate([np.asarray(y_true),
np.asarray(y_pred)]))

@classmethod
def accuracy(cls, y_true, y_pred):
    """correct predictions / all predictions"""
    y_true, y_pred = np.asarray(y_true), np.asarray(y_pred)
    return float(np.sum(y_true == y_pred) / len(y_true))

@classmethod
def precision(cls, y_true, y_pred, classes=None):
    """precision_c = TP_c / (TP_c + FP_c), one value per class"""
    out = []
    for c in cls._classes(y_true, y_pred, classes):
        tp, fp, fn = cls._confusion_counts(y_true, y_pred, c)
        out.append(tp / (tp + fp) if (tp + fp) > 0 else 0.0)
    return np.array(out)

@classmethod
def recall(cls, y_true, y_pred, classes=None):

```

```

        """recall_c = TP_c / (TP_c + FN_c), one value per class"""
        out = []
        for c in cls._classes(y_true, y_pred, classes):
            tp, fp, fn = cls._confusion_counts(y_true, y_pred, c)
            out.append(tp / (tp + fn) if (tp + fn) > 0 else 0.0)
        return np.array(out)

    @classmethod
    def f1_score(cls, y_true, y_pred, classes=None):
        """f1_c = 2 * precision_c * recall_c / (precision_c +
recall_c)"""
        p = cls.precision(y_true, y_pred, classes)
        r = cls.recall(y_true, y_pred, classes)
        denom = p + r
        return np.divide(2 * p * r, denom, out=np.zeros_like(denom),
where=denom > 0)

    @classmethod
    def support(cls, y_true, classes=None):
        """number of true samples in each class"""
        y_true = np.asarray(y_true)
        if classes is None:
            classes = np.unique(y_true)
        return np.array([int(np.sum(y_true == c)) for c in classes])

    # macro: simple average over the classes
    @classmethod
    def macro_precision(cls, y_true, y_pred, classes=None):
        return float(cls.precision(y_true, y_pred, classes).mean())

    @classmethod
    def macro_recall(cls, y_true, y_pred, classes=None):
        return float(cls.recall(y_true, y_pred, classes).mean())

    @classmethod
    def macro_f1(cls, y_true, y_pred, classes=None):
        return float(cls.f1_score(y_true, y_pred, classes).mean())

    # weighted: each class weighted by its share of the true labels
    @classmethod
    def _weighted(cls, scores, y_true, classes):
        s = cls.support(y_true, classes)
        return float(np.sum(scores * s) / np.sum(s))

    @classmethod
    def weighted_precision(cls, y_true, y_pred, classes=None):
        classes = cls._classes(y_true, y_pred, classes)
        return cls._weighted(cls.precision(y_true, y_pred, classes),
y_true, classes)

    @classmethod
    def weighted_recall(cls, y_true, y_pred, classes=None):
        classes = cls._classes(y_true, y_pred, classes)
        return cls._weighted(cls.recall(y_true, y_pred, classes), y_true,
classes)

    @classmethod
    def weighted_f1(cls, y_true, y_pred, classes=None):
        classes = cls._classes(y_true, y_pred, classes)

```

```

        return cls._weighted(cls.f1_score(y_true, y_pred, classes),
                               y_true, classes)

    @classmethod
    def classification_report(cls, y_true, y_pred, classes=None,
                              digits=2):
        """Same layout as sklearn.metrics.classification_report."""
        classes = cls._classes(y_true, y_pred, classes)
        p, r = cls.precision(y_true, y_pred, classes), cls.recall(y_true,
                                                                    y_pred, classes)
        f1, s = cls.f1_score(y_true, y_pred, classes),
        cls.support(y_true, classes)

        lines = [f"{' ':>12}{'precision':>12}{'recall':>12}{'f1-
score':>12}{'support':>12}", ""]
        for i, c in enumerate(classes):
            lines.append(f"{str(c):>12}{p[i]:>12.{digits}f}{r[i]:>12.
{digits}f}{f1[i]:>12.{digits}f}{s[i]:>12d}")
            lines.append("")
            acc, total = cls.accuracy(y_true, y_pred), int(np.sum(s))
            lines.append(f"{'accuracy':>12}{' ':>12}{' ':>12}{acc:>12.
{digits}f}{total:>12d}")
            lines.append(f"{'macro avg':>12}{p.mean():>12.{digits}f}
{r.mean():>12.{digits}f}{f1.mean():>12.{digits}f}{total:>12d}")
            wp = cls.weighted_precision(y_true, y_pred, classes)
            wr = cls.weighted_recall(y_true, y_pred, classes)
            wfl = cls.weighted_f1(y_true, y_pred, classes)
            lines.append(f"{'weighted avg':>12}{wp:>12.{digits}f}{wr:>12.
{digits}f}{wfl:>12.{digits}f}{total:>12d}")
        return "\n".join(lines)

```

3.2 About the weighted average

Weighted precision multiplies each class's precision by that class's share of the true labels (its `support` divided by the total) and adds them up. For example, with 20% / 30% / 20% / 30% of the samples in classes 0 to 3:

$$\text{weighted precision} = 0.2 p_0 + 0.3 p_1 + 0.2 p_2 + 0.3 p_3$$

The weights already add up to 1, so I do not divide by 4 again. The assignment PDF shows the weighted formula with a division by 4, but with it the result would be 4 times smaller than sklearn's `weighted avg`. Since the task asks to compare with sklearn, I follow the sklearn definition, and the numbers below match it exactly.

For the rest of the notebook I make short names for the metric methods so the code is easier to read:

```
In [11]: accuracy = LogisticRegression.accuracy
macro_precision = LogisticRegression.macro_precision
macro_recall = LogisticRegression.macro_recall
macro_f1 = LogisticRegression.macro_f1
weighted_precision = LogisticRegression.weighted_precision
weighted_recall = LogisticRegression.weighted_recall
weighted_f1 = LogisticRegression.weighted_f1
classification_report_scratch = LogisticRegression.classification_report
```

3.3 Compare with `sklearn.metrics.classification_report` (mock data)

I made the mock data **imbalanced on purpose** (50 / 30 / 15 / 5 samples per class). With balanced classes, macro and weighted averages would be almost the same, so the check would not really test the weighted part.

```
In [12]: rng = np.random.RandomState(0)
mock_y_true = np.array([0] * 50 + [1] * 30 + [2] * 15 + [3] * 5)
rng.shuffle(mock_y_true)
mock_y_pred = mock_y_true.copy()
flip_idx = rng.choice(len(mock_y_pred), size=20, replace=False)
mock_y_pred[flip_idx] = rng.randint(0, 4, size=20) # make 20
predictions random

print("=" * 26, "FROM SCRATCH", "=" * 26)
print(LogisticRegression.classification_report(mock_y_true, mock_y_pred))
print()
print("=" * 28, "SKLEARN", "=" * 28)
print(sk_classification_report(mock_y_true, mock_y_pred, digits=2))
```

```
===== FROM SCRATCH =====
              precision      recall    f1-score      support

   0              0.93          0.84          0.88          50
   1              0.96          0.83          0.89          30
   2              0.74          0.93          0.82          15
   3              0.30          0.60          0.40           5

 accuracy              0.84          100
 macro avg              0.73          0.80          0.75          100
 weighted avg           0.88          0.84          0.85          100

===== SKLEARN =====
              precision      recall    f1-score      support

   0              0.93          0.84          0.88          50
   1              0.96          0.83          0.89          30
   2              0.74          0.93          0.82          15
   3              0.30          0.60          0.40           5

 accuracy              0.84          100
 macro avg              0.73          0.80          0.75          100
 weighted avg           0.88          0.84          0.85          100
```

Every number (per class precision, recall, f1, accuracy, macro avg and weighted avg) is the same as sklearn's report.

Q: What does support mean in the classification report?

A: It is the number of *true* samples of each class in `y_true`. It does not depend on the predictions. It shows how many samples each class's score is based on (a class with support 5 gives a much less reliable score than one with support 500), and it is the weight used in the `weighted avg` row.

3.4 Train on the Car Price data (no penalty)


The results here are on the **validation** set. The test set is not used until Section 6.

```
In [13]: model_base = LogisticRegression(
          k=k, n=n, method="batch", alpha=1.0, max_iter=5000,
          use_penalty=False, verbose=True, random_state=1,
        )
          model_base.fit(X_train, Y_train_onehot)
```

```
iter 0: loss=1.9131
iter 1000: loss=0.5911
iter 2000: loss=0.5870
iter 3000: loss=0.5859
iter 4000: loss=0.5853
time taken: 2.79s
```

```
Out[13]: <__main__.LogisticRegression at 0x1aba85ba490>
```

```
In [14]: plt.figure(figsize=(6, 4))
          plt.plot(model_base.losses)
          plt.title("Training loss (no penalty)")
          plt.xlabel("iteration")
          plt.ylabel("loss")
          plt.show()
```



```
In [15]: yhat_base = model_base.predict(X_val)

print("=" * 22, "FROM SCRATCH (validation)", "=" * 22)
print(model_base.classification_report(y_val, yhat_base))  # metric
               method on the trained model
print()
print("=" * 24, "SKLEARN (validation)", "=" * 24)
print(sk_classification_report(y_val, yhat_base, digits=2))
```

```
===== FROM SCRATCH (validation) =====
              precision    recall  f1-score   support

     0           0.89         0.80         0.84         311
     1           0.61         0.64         0.62         242
     2           0.62         0.67         0.65         261
     3           0.83         0.83         0.83         276

 accuracy                   0.74         1090
 macro avg           0.74         0.74         0.74         1090
 weighted avg        0.75         0.74         0.74         1090
```

```
===== SKLEARN (validation) =====
              precision    recall  f1-score   support

     0           0.89         0.80         0.84         311
     1           0.61         0.64         0.62         242
     2           0.62         0.67         0.65         261
     3           0.83         0.83         0.83         276

 accuracy                   0.74         1090
 macro avg           0.74         0.74         0.74         1090
 weighted avg        0.75         0.74         0.74         1090
```

4. Task 2: Ridge (L2) penalty

$$J(\theta) = - \sum_{i=1}^m y^{(i)} \log(h^{(i)}) + \lambda \sum_{j=1}^n \theta_j^2$$

`RidgePenalty` is adapted from `03 - Bias-Variance Tradeoff and Regularization.ipynb`. Here θ is a weight **matrix** W (features x classes) instead of a vector, so the penalty is the sum of squares of all its entries. The bias row is not penalized. It is turned on with `LogisticRegression(..., use_penalty=True, l=<lambda>)` and off with `use_penalty=False` (the default).

Scaling in the code. The formula above is the unscaled loss. In the code the loss is averaged: the cross entropy is divided by the batch size `m` and the penalty by the training set size `N`, so the code minimizes $\frac{1}{m} \text{CE} + \frac{\lambda}{N} \sum_j \theta_j^2$. For batch gradient descent `m = N`, so this is exactly $J(\theta)/N$: the same minimum, only a different scale. The one practical effect is that a given value of λ here is not directly comparable with an unscaled implementation. Dividing the penalty by `N` (not by the batch size) keeps its strength the same for batch, minibatch and stochastic updates.

```
In [16]: model_ridge = LogisticRegression(
          k=k, n=n, method="batch", alpha=1.0, max_iter=5000,
          use_penalty=True, l=0.5, verbose=True, random_state=1,
          )
          model_ridge.fit(X_train, Y_train_onehot)
          yhat_ridge = model_ridge.predict(X_val)
```

```
iter 0: loss=1.9177
iter 1000: loss=0.6086
iter 2000: loss=0.6057
iter 3000: loss=0.6048
iter 4000: loss=0.6043
time taken: 2.71s
```

```
In [17]: # compared on VALIDATION, not test
comparison = pd.DataFrame({
    "metric": ["accuracy", "macro_precision", "macro_recall", "macro_f1",
               "weighted_precision", "weighted_recall", "weighted_f1"],
    "no_penalty": [accuracy(y_val, yhat_base), macro_precision(y_val,
                                                                yhat_base),
                   macro_recall(y_val, yhat_base), macro_f1(y_val,
                                                             yhat_base),
                   weighted_precision(y_val, yhat_base),
                   weighted_recall(y_val, yhat_base),
                   weighted_f1(y_val, yhat_base)],
    "ridge_l2 (l=0.5)": [accuracy(y_val, yhat_ridge),
                         macro_precision(y_val, yhat_ridge),
                         macro_recall(y_val, yhat_ridge), macro_f1(y_val,
                                                                     yhat_ridge),
                         weighted_precision(y_val, yhat_ridge),
                         weighted_recall(y_val, yhat_ridge),
                         weighted_f1(y_val, yhat_ridge)],
})
comparison
```

```
Out[17]:
```

	metric	no_penalty	ridge_l2 (l=0.5)
0	accuracy	0.741284	0.744037
1	macro_precision	0.738374	0.739996
2	macro_recall	0.735210	0.737923
3	macro_f1	0.735926	0.738211
4	weighted_precision	0.748781	0.750326
5	weighted_recall	0.741284	0.744037
6	weighted_f1	0.744101	0.746371

The penalty should make the weights smaller. Sum of squared weights (without the bias row):

```
In [18]: print("no penalty    -- sum(W[1:]**2):",
              round(float(np.sum(model_base.W[1:, :] ** 2)), 3))
print("ridge (l=0.5) -- sum(W[1:]**2):",
              round(float(np.sum(model_ridge.W[1:, :] ** 2)), 3))
```

no penalty -- sum(W[1:]**2): 221.477
ridge (l=0.5) -- sum(W[1:]**2): 138.893

5. Task 3.1: Log the experiments with MLflow

Experiment name `<student_ID>-a3`. Params, metrics and the model are logged, the dataset is not.

Note: the course MLflow server (`mlflow.ml.brain.cs.ait.ac.th`) was down. Following the TA announcement on 1 Oct 2026, logging to it is no longer required, and screenshots of a local MLflow instance are submitted instead (see the README). The cell below still tries the course server first. If it cannot connect, it logs to a local SQLite store (`mlflow_fallback.db`) like in A2.

```

In [19]: import mlflow
import mlflow.pyfunc
from mlflow.tracking import MlflowClient

AIT_MLFLOW_URI = "http://mlflow.ml.brain.cs.ait.ac.th/"
EXPERIMENT_NAME = f"{STUDENT_ID}-a3"

try:
    mlflow.set_tracking_uri(AIT_MLFLOW_URI)
    mlflow.set_experiment(EXPERIMENT_NAME)
    _ = MlflowClient().search_experiments() # makes a real request to
the server
    print(f"Connected to AIT MLflow server, logging to '{EXPERIMENT_NAME}'")
except Exception as e:
    fallback_uri = "sqlite:///mlflow_fallback.db"
    mlflow.set_tracking_uri(fallback_uri)
    mlflow.set_experiment(EXPERIMENT_NAME)
    print(f"Could not reach {AIT_MLFLOW_URI} ({type(e).__name__})")
    print(f"Logging to the local tracking store '{fallback_uri}' instead.")

```

```
2026/10/02 17:24:46 INFO mlflow.store.db.utils: Creating initial MLflow
database tables...
2026/10/02 17:24:46 INFO mlflow.store.db.utils: Updating database tables
INFO [alembic.runtime.migration] Context impl SQLiteImpl.
INFO [alembic.runtime.migration] Will assume non-transactional DDL.
INFO [alembic.runtime.migration] Running upgrade -> 451aebb31d03, add
metric step
INFO [alembic.runtime.migration] Running upgrade 451aebb31d03 ->
90e64c465722, migrate user column to tags
INFO [alembic.runtime.migration] Running upgrade 90e64c465722 ->
181f10493468, allow nulls for metric values
INFO [alembic.runtime.migration] Running upgrade 181f10493468 ->
df50e92ffc5e, Add Experiment Tags Table
INFO [alembic.runtime.migration] Running upgrade df50e92ffc5e ->
7ac759974ad8, Update run tags with larger limit
INFO [alembic.runtime.migration] Running upgrade 7ac759974ad8 ->
89d4b8295536, create latest metrics table
INFO [89d4b8295536_create_latest_metrics_table_py] Migration complete!
INFO [alembic.runtime.migration] Running upgrade 89d4b8295536 ->
2b4d017a5e9b, add model registry tables to db
INFO [2b4d017a5e9b_add_model_registry_tables_to_db_py] Adding
registered_models and model_versions tables to database.
INFO [2b4d017a5e9b_add_model_registry_tables_to_db_py] Migration
complete!
INFO [alembic.runtime.migration] Running upgrade 2b4d017a5e9b ->
cfd24bdc0731, Update run status constraint with killed
INFO [alembic.runtime.migration] Running upgrade cfd24bdc0731 ->
0a8213491aaa, drop_duplicate_killed_constraint
INFO [alembic.runtime.migration] Running upgrade 0a8213491aaa ->
728d730b5ebd, add registered model tags table
INFO [alembic.runtime.migration] Running upgrade 728d730b5ebd ->
27a6a02d2cf1, add model version tags table
INFO [alembic.runtime.migration] Running upgrade 27a6a02d2cf1 ->
84291f40a231, add run_link to model_version
INFO [alembic.runtime.migration] Running upgrade 84291f40a231 ->
a8c4a736bde6, allow nulls for run_id
INFO [alembic.runtime.migration] Running upgrade a8c4a736bde6 ->
39d1c3be5f05, add_is_nan_constraint_for_metrics_tables_if_necessary
INFO [alembic.runtime.migration] Running upgrade 39d1c3be5f05 ->
c48cb773bb87, reset_default_value_for_is_nan_in_metrics_table_for_mysql
INFO [alembic.runtime.migration] Running upgrade c48cb773bb87 ->
bd07f7e963c5, create index on run_uuid
INFO [alembic.runtime.migration] Running upgrade bd07f7e963c5 ->
0c779009ac13, add deleted_time field to runs table
INFO [alembic.runtime.migration] Running upgrade 0c779009ac13 ->
cc1f77228345, change param value length to 500
INFO [alembic.runtime.migration] Running upgrade cc1f77228345 ->
97727af70f4d, Add creation_time and last_update_time to experiments table
INFO [alembic.runtime.migration] Running upgrade 97727af70f4d ->
3500859a5d39, Add Model Aliases table
INFO [alembic.runtime.migration] Running upgrade 3500859a5d39 ->
7f2a7d5fae7d, add datasets inputs input_tags tables
INFO [alembic.runtime.migration] Running upgrade 7f2a7d5fae7d ->
2d6e25af4d3e, increase max param val length from 500 to 8000
INFO [alembic.runtime.migration] Running upgrade 2d6e25af4d3e ->
acf3f17fdcc7, add storage location field to model versions
INFO [alembic.runtime.migration] Running upgrade acf3f17fdcc7 ->
867495a8f9d4, add trace tables
INFO [alembic.runtime.migration] Running upgrade 867495a8f9d4 ->
```

```

5b0e9adcef9c, add cascade deletion to trace tables foreign keys
INFO [alembic.runtime.migration] Running upgrade 5b0e9adcef9c ->
4465047574b1, increase max dataset schema size
INFO [alembic.runtime.migration] Running upgrade 4465047574b1 ->
f5a4f2784254, increase run tag value limit to 8000
INFO [alembic.runtime.migration] Running upgrade f5a4f2784254 ->
0584bdc529eb, add cascading deletion to datasets from experiments
INFO [alembic.runtime.migration] Running upgrade 0584bdc529eb ->
400f98739977, add logged model tables
INFO [alembic.runtime.migration] Running upgrade 400f98739977 ->
6953534de441, add step to inputs table
INFO [alembic.runtime.migration] Running upgrade 6953534de441 ->
bda7b8c39065, increase_model_version_tag_value_limit
INFO [alembic.runtime.migration] Context impl SQLiteImpl.
INFO [alembic.runtime.migration] Will assume non-transactional DDL.
2026/10/02 17:24:47 INFO mlflow.tracking.fluent: Experiment with name
'st127031-a3' does not exist. Creating a new experiment.
Could not reach http://mlflow.ml.brain.cs.ait.ac.th/ (MlflowException)
Logging to the local tracking store 'sqlite:///mlflow_fallback.db'
instead.

```

A small grid of experiments (gradient descent method, learning rate, ridge on/off). Each one is its own MLflow run with params and **validation** metrics.

The model is saved as an MLflow **pyfunc** that only needs two objects to predict: the fitted `prep` (`ColumnTransformer`) and `W` (a numpy array), not the `LogisticRegression` object itself. This is the same idea as A2, where only `theta` and `prep` were exported. The web app then does not need this notebook's class to load the model.

Note on `final_train_loss` : for the runs with `use_penalty=True` this value includes the ridge penalty, so it is not comparable with the runs without penalty (a higher value does not mean the model fits worse). The runs are compared only with the validation metrics.

```

In [20]: class CarPriceClassifier(mlflow.pyfunc.PythonModel):
    """Prediction = prep.transform(X), add bias column, X @ W, softmax,
    argmax."""

    def load_context(self, context):
        import joblib
        bundle = joblib.load(context.artifacts["bundle"])
        self.prep = bundle["prep"]
        self.W = bundle["W"]
        self.feature_cols = bundle["feature_cols"]
        self.class_labels = bundle["class_labels"]
        self.common_brands = bundle["common_brands"]

    def _group_brand(self, b):
        """Same brand grouping as in training. Matching ignores
upper/lower
case, so "maruti" in the web app is treated as "Maruti"."""
        lookup = {c.lower(): c for c in self.common_brands}
        return lookup.get(str(b).strip().lower(), "Other")

    def _softmax(self, z):
        z = z - np.max(z, axis=1, keepdims=True)
        e = np.exp(z)
        return e / np.sum(e, axis=1, keepdims=True)

    def predict(self, context, model_input):
        model_input = model_input[self.feature_cols].copy()
        model_input["brand"] =
model_input["brand"].apply(self._group_brand)
        X = self.prep.transform(model_input)
        if hasattr(X, "toarray"):
            X = X.toarray()
        X = np.concatenate([np.ones((X.shape[0], 1)), X], axis=1)
        probs = self._softmax(X @ self.W)
        preds = np.argmax(probs, axis=1)
        out = pd.DataFrame({"predicted_class": preds})
        for i in range(probs.shape[1]):
            out[f"proba_class_{i}"] = probs[:, i]
        return out

import os
import joblib

os.makedirs("dev/_tmp", exist_ok=True) # temporary model files for
MLflow (ignored by git)

# no penalty with two learning rates and minibatch, plus ridge with a
range of
# lambda values from very weak (0.01) to very strong (10)
experiment_grid = [
    dict(method="batch", alpha=0.5, max_iter=5000, use_penalty=False,
l=0.0),
    dict(method="batch", alpha=1.0, max_iter=5000, use_penalty=False,
l=0.0),
    dict(method="minibatch", alpha=2.0, max_iter=3000, use_penalty=False,
l=0.0),
] + [
    dict(method="batch", alpha=1.0, max_iter=5000, use_penalty=True,

```

```

l=lam)
    for lam in [0.01, 0.1, 0.5, 1.0, 5.0, 10.0]
]

run_results = []
trained_models = {} # run_id -> fitted model, so the best one does not
need retraining

for cfg in experiment_grid:
    run_name = f"{cfg['method']}-alpha{cfg['alpha']}-
penalty{cfg['use_penalty']}-l{cfg['l']}"
    with mlflow.start_run(run_name=run_name) as run:
        mlflow.log_params(**cfg, "k": k, "n": n, "student_id":
STUDENT_ID})

        mdl = LogisticRegression(k=k, n=n, random_state=1, **cfg)
        mdl.fit(X_train, Y_train_onehot)
        yhat = mdl.predict(X_val) # model selection uses VALIDATION only

        metrics = {
            "val_accuracy": mdl.accuracy(y_val, yhat),
            "val_macro_precision": mdl.macro_precision(y_val, yhat),
            "val_macro_recall": mdl.macro_recall(y_val, yhat),
            "val_macro_f1": mdl.macro_f1(y_val, yhat),
            "val_weighted_precision": mdl.weighted_precision(y_val,
yhat),
            "val_weighted_recall": mdl.weighted_recall(y_val, yhat),
            "val_weighted_f1": mdl.weighted_f1(y_val, yhat),
            "final_train_loss": mdl.losses[-1],
        }
        mlflow.log_metrics(metrics)
        for step, loss_val in enumerate(mdl.losses):
            if step % 200 == 0:
                mlflow.log_metric("train_loss_curve", loss_val,
step=step)

        bundle_path = f"dev/_tmp/bundle_{run_name}.pkl"
        joblib.dump({
            "W": mdl.W, "prep": prep, "feature_cols": FEATURE_COLS,
            "common_brands": COMMON_BRANDS,
            "class_labels": CLASS_LABELS, "class_price_ranges":
CLASS_PRICE_RANGES,
        }, bundle_path)
        model_info = mlflow.pyfunc.log_model(
            name="model",
            python_model=CarPriceClassifier(),
            artifacts={"bundle": bundle_path},
            input_example=df_train[FEATURE_COLS].dropna().head(3), #
gives the model a signature
        )

        trained_models[run.info.run_id] = mdl
        run_results.append({"run_id": run.info.run_id, "run_name":
run_name,
                                "model_uri": model_info.model_uri, **cfg,
**metrics})
        print(f"[{run_name}] val_accuracy={metrics['val_accuracy']:.4f}
val_macro_f1={metrics['val_macro_f1']:.4f}")

```

```
results_df = pd.DataFrame(run_results).sort_values("val_macro_f1",
ascending=False)
results_df
```

```
c:\Users\ASUS\.conda\envs\ML\lib\site-
packages\mlflow\pyfunc\utils\data_validation.py:186: UserWarning: Add type
hints to the `predict` method to enable data validation and automatic
signature inference during model logging. Check
https://mlflow.org/docs/latest/model/python\_model.html#type-hint-usage-in-
pythonmodel for more details.
  color_warning(
2026/10/02 17:24:51 INFO mlflow.pyfunc: Inferring model signature from
input example
c:\Users\ASUS\.conda\envs\ML\lib\site-packages\mlflow\types\utils.py:452:
UserWarning: Hint: Inferred schema contains integer column(s). Integer
columns in Python cannot represent missing values. If your input data
contains missing values at inference time, it will be encoded as floats
and will cause a schema enforcement error. The best way to avoid this
problem is to infer the model schema based on a realistic data sample
(training dataset) that includes missing values. Alternatively, you can
declare integer columns as doubles (float64) whenever these columns may
have missing values. See `Handling Integers With Missing Values
<https://www.mlflow.org/docs/latest/models.html#handling-integers-with-
missing-values>`_ for more details.
  warnings.warn(
c:\Users\ASUS\.conda\envs\ML\lib\site-packages\mlflow\types\utils.py:452:
UserWarning: Hint: Inferred schema contains integer column(s). Integer
columns in Python cannot represent missing values. If your input data
contains missing values at inference time, it will be encoded as floats
and will cause a schema enforcement error. The best way to avoid this
problem is to infer the model schema based on a realistic data sample
(training dataset) that includes missing values. Alternatively, you can
declare integer columns as doubles (float64) whenever these columns may
have missing values. See `Handling Integers With Missing Values
<https://www.mlflow.org/docs/latest/models.html#handling-integers-with-
missing-values>`_ for more details.
  warnings.warn(
c:\Users\ASUS\.conda\envs\ML\lib\site-packages\tqdm\auto.py:21:
TqdmWarning: IProgress not found. Please update jupyter and ipywidgets.
See https://ipywidgets.readthedocs.io/en/stable/user\_install.html
  from .autonotebook import tqdm as notebook_tqdm
Downloading artifacts: 100%|██████████| 1/1 [00:00<00:00, 1000.55it/s]
Downloading artifacts: 100%|██████████| 8/8 [00:00<00:00, 1333.54it/s]
[batch-alpha0.5-penaltyFalse-l0.0] val_accuracy=0.7385
val_macro_f1=0.7332
```

```

2026/10/02 17:25:04 INFO mlflow.pyfunc: Inferring model signature from
input example
c:\Users\ASUS\.conda\envs\ML\lib\site-packages\mlflow\types\utils.py:452:
UserWarning: Hint: Inferred schema contains integer column(s). Integer
columns in Python cannot represent missing values. If your input data
contains missing values at inference time, it will be encoded as floats
and will cause a schema enforcement error. The best way to avoid this
problem is to infer the model schema based on a realistic data sample
(training dataset) that includes missing values. Alternatively, you can
declare integer columns as doubles (float64) whenever these columns may
have missing values. See `Handling Integers With Missing Values
<https://www.mlflow.org/docs/latest/models.html#handling-integers-with-
missing-values>`_ for more details.
  warnings.warn(
c:\Users\ASUS\.conda\envs\ML\lib\site-packages\mlflow\types\utils.py:452:
UserWarning: Hint: Inferred schema contains integer column(s). Integer
columns in Python cannot represent missing values. If your input data
contains missing values at inference time, it will be encoded as floats
and will cause a schema enforcement error. The best way to avoid this
problem is to infer the model schema based on a realistic data sample
(training dataset) that includes missing values. Alternatively, you can
declare integer columns as doubles (float64) whenever these columns may
have missing values. See `Handling Integers With Missing Values
<https://www.mlflow.org/docs/latest/models.html#handling-integers-with-
missing-values>`_ for more details.
  warnings.warn(
Downloading artifacts: 100%|██████████| 1/1 [00:00<00:00, 334.85it/s]
Downloading artifacts: 100%|██████████| 8/8 [00:00<00:00, 1052.13it/s]
[batch-alpha1.0-penaltyFalse-l0.0] val_accuracy=0.7413
val_macro_f1=0.7359

```

```

2026/10/02 17:25:09 INFO mlflow.pyfunc: Inferring model signature from
input example
c:\Users\ASUS\.conda\envs\ML\lib\site-packages\mlflow\types\utils.py:452:
UserWarning: Hint: Inferred schema contains integer column(s). Integer
columns in Python cannot represent missing values. If your input data
contains missing values at inference time, it will be encoded as floats
and will cause a schema enforcement error. The best way to avoid this
problem is to infer the model schema based on a realistic data sample
(training dataset) that includes missing values. Alternatively, you can
declare integer columns as doubles (float64) whenever these columns may
have missing values. See `Handling Integers With Missing Values
<https://www.mlflow.org/docs/latest/models.html#handling-integers-with-
missing-values>`_ for more details.
  warnings.warn(
c:\Users\ASUS\.conda\envs\ML\lib\site-packages\mlflow\types\utils.py:452:
UserWarning: Hint: Inferred schema contains integer column(s). Integer
columns in Python cannot represent missing values. If your input data
contains missing values at inference time, it will be encoded as floats
and will cause a schema enforcement error. The best way to avoid this
problem is to infer the model schema based on a realistic data sample
(training dataset) that includes missing values. Alternatively, you can
declare integer columns as doubles (float64) whenever these columns may
have missing values. See `Handling Integers With Missing Values
<https://www.mlflow.org/docs/latest/models.html#handling-integers-with-
missing-values>`_ for more details.
  warnings.warn(
Downloading artifacts: 100%|██████████| 1/1 [00:00<00:00, 664.39it/s]
Downloading artifacts: 100%|██████████| 8/8 [00:00<00:00, 1449.56it/s]

```

```
[minibatch-alpha2.0-penaltyFalse-l0.0] val_accuracy=0.7422  
val_macro_f1=0.7350
```

```
2026/10/02 17:25:17 INFO mlflow.pyfunc: Inferring model signature from  
input example
```

```
c:\Users\ASUS\.conda\envs\ML\lib\site-packages\mlflow\types\utils.py:452:  
UserWarning: Hint: Inferred schema contains integer column(s). Integer  
columns in Python cannot represent missing values. If your input data  
contains missing values at inference time, it will be encoded as floats  
and will cause a schema enforcement error. The best way to avoid this  
problem is to infer the model schema based on a realistic data sample  
(training dataset) that includes missing values. Alternatively, you can  
declare integer columns as doubles (float64) whenever these columns may  
have missing values. See `Handling Integers With Missing Values`  
<https://www.mlflow.org/docs/latest/models.html#handling-integers-with-  
missing-values>`_ for more details.
```

```
warnings.warn(  
c:\Users\ASUS\.conda\envs\ML\lib\site-packages\mlflow\types\utils.py:452:  
UserWarning: Hint: Inferred schema contains integer column(s). Integer  
columns in Python cannot represent missing values. If your input data  
contains missing values at inference time, it will be encoded as floats  
and will cause a schema enforcement error. The best way to avoid this  
problem is to infer the model schema based on a realistic data sample  
(training dataset) that includes missing values. Alternatively, you can  
declare integer columns as doubles (float64) whenever these columns may  
have missing values. See `Handling Integers With Missing Values`  
<https://www.mlflow.org/docs/latest/models.html#handling-integers-with-  
missing-values>`_ for more details.
```

```
warnings.warn(  
Downloading artifacts: 100%|██████████| 1/1 [00:00<00:00, 499.02it/s]  
Downloading artifacts: 100%|██████████| 8/8 [00:00<00:00, 1998.72it/s]
```

```
[batch-alpha1.0-penaltyTrue-l0.01] val_accuracy=0.7404  
val_macro_f1=0.7349
```

```

2026/10/02 17:25:26 INFO mlflow.pyfunc: Inferring model signature from
input example
c:\Users\ASUS\.conda\envs\ML\lib\site-packages\mlflow\types\utils.py:452:
UserWarning: Hint: Inferred schema contains integer column(s). Integer
columns in Python cannot represent missing values. If your input data
contains missing values at inference time, it will be encoded as floats
and will cause a schema enforcement error. The best way to avoid this
problem is to infer the model schema based on a realistic data sample
(training dataset) that includes missing values. Alternatively, you can
declare integer columns as doubles (float64) whenever these columns may
have missing values. See `Handling Integers With Missing Values
<https://www.mlflow.org/docs/latest/models.html#handling-integers-with-
missing-values>`_ for more details.
  warnings.warn(
c:\Users\ASUS\.conda\envs\ML\lib\site-packages\mlflow\types\utils.py:452:
UserWarning: Hint: Inferred schema contains integer column(s). Integer
columns in Python cannot represent missing values. If your input data
contains missing values at inference time, it will be encoded as floats
and will cause a schema enforcement error. The best way to avoid this
problem is to infer the model schema based on a realistic data sample
(training dataset) that includes missing values. Alternatively, you can
declare integer columns as doubles (float64) whenever these columns may
have missing values. See `Handling Integers With Missing Values
<https://www.mlflow.org/docs/latest/models.html#handling-integers-with-
missing-values>`_ for more details.
  warnings.warn(
Downloading artifacts: 100%|██████████| 1/1 [00:00<00:00, 1000.07it/s]
Downloading artifacts: 100%|██████████| 8/8 [00:00<00:00, 1598.36it/s]
[batch-alpha1.0-penaltyTrue-l0.1] val_accuracy=0.7385 val_macro_f1=0.7331

2026/10/02 17:25:34 INFO mlflow.pyfunc: Inferring model signature from
input example
c:\Users\ASUS\.conda\envs\ML\lib\site-packages\mlflow\types\utils.py:452:
UserWarning: Hint: Inferred schema contains integer column(s). Integer
columns in Python cannot represent missing values. If your input data
contains missing values at inference time, it will be encoded as floats
and will cause a schema enforcement error. The best way to avoid this
problem is to infer the model schema based on a realistic data sample
(training dataset) that includes missing values. Alternatively, you can
declare integer columns as doubles (float64) whenever these columns may
have missing values. See `Handling Integers With Missing Values
<https://www.mlflow.org/docs/latest/models.html#handling-integers-with-
missing-values>`_ for more details.
  warnings.warn(
c:\Users\ASUS\.conda\envs\ML\lib\site-packages\mlflow\types\utils.py:452:
UserWarning: Hint: Inferred schema contains integer column(s). Integer
columns in Python cannot represent missing values. If your input data
contains missing values at inference time, it will be encoded as floats
and will cause a schema enforcement error. The best way to avoid this
problem is to infer the model schema based on a realistic data sample
(training dataset) that includes missing values. Alternatively, you can
declare integer columns as doubles (float64) whenever these columns may
have missing values. See `Handling Integers With Missing Values
<https://www.mlflow.org/docs/latest/models.html#handling-integers-with-
missing-values>`_ for more details.
  warnings.warn(
Downloading artifacts: 100%|██████████| 1/1 [00:00<00:00, 993.20it/s]
Downloading artifacts: 100%|██████████| 8/8 [00:00<00:00, 1753.29it/s]
[batch-alpha1.0-penaltyTrue-l0.5] val_accuracy=0.7440 val_macro_f1=0.7382

```

```

2026/10/02 17:25:43 INFO mlflow.pyfunc: Inferring model signature from
input example
c:\Users\ASUS\.conda\envs\ML\lib\site-packages\mlflow\types\utils.py:452:
UserWarning: Hint: Inferred schema contains integer column(s). Integer
columns in Python cannot represent missing values. If your input data
contains missing values at inference time, it will be encoded as floats
and will cause a schema enforcement error. The best way to avoid this
problem is to infer the model schema based on a realistic data sample
(training dataset) that includes missing values. Alternatively, you can
declare integer columns as doubles (float64) whenever these columns may
have missing values. See `Handling Integers With Missing Values
<https://www.mlflow.org/docs/latest/models.html#handling-integers-with-
missing-values>`_ for more details.
  warnings.warn(
c:\Users\ASUS\.conda\envs\ML\lib\site-packages\mlflow\types\utils.py:452:
UserWarning: Hint: Inferred schema contains integer column(s). Integer
columns in Python cannot represent missing values. If your input data
contains missing values at inference time, it will be encoded as floats
and will cause a schema enforcement error. The best way to avoid this
problem is to infer the model schema based on a realistic data sample
(training dataset) that includes missing values. Alternatively, you can
declare integer columns as doubles (float64) whenever these columns may
have missing values. See `Handling Integers With Missing Values
<https://www.mlflow.org/docs/latest/models.html#handling-integers-with-
missing-values>`_ for more details.
  warnings.warn(
Downloading artifacts: 100%|██████████| 1/1 [00:00<00:00, 596.21it/s]
Downloading artifacts: 100%|██████████| 8/8 [00:00<00:00, 1333.17it/s]
[batch-alpha1.0-penaltyTrue-l1.0] val_accuracy=0.7404 val_macro_f1=0.7340

2026/10/02 17:25:50 INFO mlflow.pyfunc: Inferring model signature from
input example
c:\Users\ASUS\.conda\envs\ML\lib\site-packages\mlflow\types\utils.py:452:
UserWarning: Hint: Inferred schema contains integer column(s). Integer
columns in Python cannot represent missing values. If your input data
contains missing values at inference time, it will be encoded as floats
and will cause a schema enforcement error. The best way to avoid this
problem is to infer the model schema based on a realistic data sample
(training dataset) that includes missing values. Alternatively, you can
declare integer columns as doubles (float64) whenever these columns may
have missing values. See `Handling Integers With Missing Values
<https://www.mlflow.org/docs/latest/models.html#handling-integers-with-
missing-values>`_ for more details.
  warnings.warn(
c:\Users\ASUS\.conda\envs\ML\lib\site-packages\mlflow\types\utils.py:452:
UserWarning: Hint: Inferred schema contains integer column(s). Integer
columns in Python cannot represent missing values. If your input data
contains missing values at inference time, it will be encoded as floats
and will cause a schema enforcement error. The best way to avoid this
problem is to infer the model schema based on a realistic data sample
(training dataset) that includes missing values. Alternatively, you can
declare integer columns as doubles (float64) whenever these columns may
have missing values. See `Handling Integers With Missing Values
<https://www.mlflow.org/docs/latest/models.html#handling-integers-with-
missing-values>`_ for more details.
  warnings.warn(
Downloading artifacts: 100%|██████████| 1/1 [00:00<00:00, 999.83it/s]
Downloading artifacts: 100%|██████████| 8/8 [00:00<00:00, 1984.18it/s]
[batch-alpha1.0-penaltyTrue-l5.0] val_accuracy=0.7294 val_macro_f1=0.7230

```

```

2026/10/02 17:25:58 INFO mlflow.pyfunc: Inferring model signature from
input example
c:\Users\ASUS\.conda\envs\ML\lib\site-packages\mlflow\types\utils.py:452:
UserWarning: Hint: Inferred schema contains integer column(s). Integer
columns in Python cannot represent missing values. If your input data
contains missing values at inference time, it will be encoded as floats
and will cause a schema enforcement error. The best way to avoid this
problem is to infer the model schema based on a realistic data sample
(training dataset) that includes missing values. Alternatively, you can
declare integer columns as doubles (float64) whenever these columns may
have missing values. See `Handling Integers With Missing Values
<https://www.mlflow.org/docs/latest/models.html#handling-integers-with-
missing-values>`_ for more details.
  warnings.warn(
c:\Users\ASUS\.conda\envs\ML\lib\site-packages\mlflow\types\utils.py:452:
UserWarning: Hint: Inferred schema contains integer column(s). Integer
columns in Python cannot represent missing values. If your input data
contains missing values at inference time, it will be encoded as floats
and will cause a schema enforcement error. The best way to avoid this
problem is to infer the model schema based on a realistic data sample
(training dataset) that includes missing values. Alternatively, you can
declare integer columns as doubles (float64) whenever these columns may
have missing values. See `Handling Integers With Missing Values
<https://www.mlflow.org/docs/latest/models.html#handling-integers-with-
missing-values>`_ for more details.
  warnings.warn(
Downloading artifacts: 100%|██████████| 1/1 [00:00<00:00, 500.87it/s]
Downloading artifacts: 100%|██████████| 8/8 [00:00<00:00, 1772.83it/s]
[batch-alpha1.0-penaltyTrue-l10.0] val_accuracy=0.7239
val_macro_f1=0.7167

```

Out[20]:

	run_id	run_name	model_uri	method	alpha	max_iter	use_penalty	l	val_accuracy	val_macro_pr ecision	val_macro_rec all	val_macro_f1	val_weighted_ precision	val_weighted_ recall	val_weighted_ f1
5	21ea8fa499754 05ea61eeff31c b84a7	batch-alpha1.0- penaltyTrue-l0.5	models/m-f2f9 4c7a7cae434d 816136771b5a cc76	batch	1.0	5000	True	0.50	0.744037	0.739996	0.737923	0.738211	0.750326	0.744037	0.746371
1	b3dda080f92b4 b8e823fe74aec 6488a7	batch-alpha1.0- penaltyfalse-l0.0	models/m-497 c76ea515e4b8 28b8db730a270 ba94	batch	1.0	5000	False	0.00	0.741284	0.738374	0.735210	0.735926	0.748781	0.741284	0.744101
2	ddc7f3fe09d74 bd09691dfee72 3a15c9	minibatch-alph a2.0-penaltyfal se-l0.0	models/m-39e 770b28c014a3 6ab76426b24f1 77a7	minibatch	2.0	3000	False	0.00	0.742202	0.735917	0.735389	0.735025	0.745994	0.742202	0.743419
3	132a029f9b4a4 3109c920125a 1a11863	batch-alpha1.0- penaltyTrue-l0.01	models/m-259 ec55727c1483 195600c295f85 6834	batch	1.0	5000	True	0.01	0.740367	0.737291	0.734252	0.734925	0.747702	0.740367	0.743121
6	3757e3080bea 4faa1244ee96 d2b13e1	batch-alpha1.0- penaltyTrue-l1.0	models/m-b21 aa697bd15442 68e3086f9a0ee 9d1e	batch	1.0	5000	True	1.00	0.740367	0.735255	0.734017	0.734037	0.745219	0.740367	0.742143
0	1a71c45ff53a4 83d9c7918871 e676bfd	batch-alpha0.5- penaltyfalse-l0.0	models/m-55a b51c14863436 5878ec7d3118 89570	batch	0.5	5000	False	0.00	0.738532	0.735575	0.732463	0.733172	0.746190	0.738532	0.741452
4	9e87e725e0c64 0d0a690efaa16 dae4f1	batch-alpha1.0- penaltyTrue-l0.1	models/m-2f6d 0746a29348a9 8d25877eba89 2831	batch	1.0	5000	True	0.10	0.738532	0.735593	0.732388	0.733139	0.746151	0.738532	0.741426
7	d5605acc4b84 44aa914925b4 642a4763	batch-alpha1.0- penaltyTrue-l5.0	models/m-0ac 52fbdes4be4b2b 9bc7e60a877df b48	batch	1.0	5000	True	5.00	0.729358	0.724398	0.722405	0.722959	0.734296	0.729358	0.731346
8	d36d65e73bbb 408cae8fa750f de58e5c	batch-alpha1.0- penaltyTrue-l1.0	models/m-496 86305e031474 e9e7179300e9 486c5	batch	1.0	5000	True	10.00	0.723853	0.717715	0.716232	0.716720	0.727217	0.723853	0.725265

MLflow UI: experiment st127031-a3 . The 9 runs, sorted by **val_macro_f1** . Each run has its params, validation metrics and the saved model (the dataset is not logged).

The screenshot shows the MLflow UI for experiment 'st127031-a3'. The interface includes tabs for 'Runs', 'Models', 'Experimental', 'Evaluation', and 'Traces'. A search bar contains the query 'metrics.rmse < 1 and params.model = "tree"'. The runs are sorted by 'val_macro_f1' in descending order. The table below represents the data shown in the UI.

	Run Name	Created	Models	Metrics			Parameters			
				val_accuracy	val_macro_f1	val_weighted_f1	alpha	l	method	use_penalty
☐	batch-alpha1.0-penaltyTr...	3 minutes ago	model	0.744036697...	0.738210652...	0.746371484...	1.0	0.5	batch	True
☐	batch-alpha1.0-penaltyFal...	3 minutes ago	model	0.741284403...	0.735926058...	0.744100942...	1.0	0.0	batch	False
☐	minibatch-alpha2.0-penal...	3 minutes ago	model	0.742201834...	0.735025077...	0.743419137...	2.0	0.0	minibatch	False
☐	batch-alpha1.0-penaltyTr...	3 minutes ago	model	0.740366972...	0.734924879...	0.743121352...	1.0	0.01	batch	True
☐	batch-alpha1.0-penaltyTr...	3 minutes ago	model	0.740366972...	0.734036796...	0.742143225...	1.0	1.0	batch	True
☐	batch-alpha0.5-penaltyFal...	4 minutes ago	model	0.738532110...	0.733171936...	0.741452312...	0.5	0.0	batch	False
☐	batch-alpha1.0-penaltyTr...	3 minutes ago	model	0.738532110...	0.733139460...	0.741425996...	1.0	0.1	batch	True
☐	batch-alpha1.0-penaltyTr...	3 minutes ago	model	0.729357798...	0.722958682...	0.731346246...	1.0	5.0	batch	True
☐	batch-alpha1.0-penaltyTr...	2 minutes ago	model	0.723853211...	0.716719662...	0.725265131...	1.0	10.0	batch	True

6. Pick the best model (validation) and test it once

The best run is chosen only by **validation macro_f1** . After that, this one model is evaluated on the test set. This is the first and only time the test set is used. The test metrics are added to the same MLflow run.

```

In [21]: best_run = results_df.iloc[0]
best_run_id = best_run["run_id"]
best_model = trained_models[best_run_id]
print(f"Best run (by validation): {best_run['run_name']} (val_macro_f1=
{best_run['val_macro_f1']:.4f})")

yhat_test = best_model.predict(X_test)  # the ONLY place X_test / y_test
are used for scoring

print()
print("=" * 22, "FROM SCRATCH (TEST, final)", "=" * 21)
print(best_model.classification_report(y_test, yhat_test))
print()
print("=" * 24, "SKLEARN (TEST, final)", "=" * 23)
print(sk_classification_report(y_test, yhat_test, digits=2))

test_metrics = {
    "test_accuracy": best_model.accuracy(y_test, yhat_test),
    "test_macro_precision": best_model.macro_precision(y_test,
yhat_test),
    "test_macro_recall": best_model.macro_recall(y_test, yhat_test),
    "test_macro_f1": best_model.macro_f1(y_test, yhat_test),
    "test_weighted_precision": best_model.weighted_precision(y_test,
yhat_test),
    "test_weighted_recall": best_model.weighted_recall(y_test,
yhat_test),
    "test_weighted_f1": best_model.weighted_f1(y_test, yhat_test),
}
with mlflow.start_run(run_id=best_run_id):
    mlflow.log_metrics(test_metrics)

pd.DataFrame({
    "validation": [best_run["val_accuracy"], best_run["val_macro_f1"],
best_run["val_weighted_f1"]],
    "test": [test_metrics["test_accuracy"],
test_metrics["test_macro_f1"], test_metrics["test_weighted_f1"]],
}, index=["accuracy", "macro_f1", "weighted_f1"])

```

Best run (by validation): batch-alpha1.0-penaltyTrue-l0.5
(val_macro_f1=0.7382)

```
===== FROM SCRATCH (TEST, final) =====
              precision      recall    f1-score      support

      0          0.86          0.82          0.84          360
      1          0.58          0.63          0.60          289
      2          0.63          0.65          0.64          357
      3          0.84          0.81          0.82          357

    accuracy                    0.73          1363
  macro avg          0.73          0.72          0.73          1363
weighted avg          0.74          0.73          0.73          1363
```

```
===== SKLEARN (TEST, final) =====
              precision      recall    f1-score      support

      0          0.86          0.82          0.84          360
      1          0.58          0.63          0.60          289
      2          0.63          0.65          0.64          357
      3          0.84          0.81          0.82          357

    accuracy                    0.73          1363
  macro avg          0.73          0.72          0.73          1363
weighted avg          0.74          0.73          0.73          1363
```

```
Out[21]:
```

	validation	test
accuracy	0.744037	0.730007
macro_f1	0.738211	0.726217
weighted_f1	0.746371	0.732698

The validation and test scores are close (about 0.74 vs 0.73). If the test score were much lower, the model would be overfitting the model selection. If it were much higher, that could mean there is still some leakage. Neither happens here.

MLflow UI: the best run. It has the `val_*` metrics like every run, plus the `test_*` metrics that were computed once on the chosen model.

st127031-a3 >

batch-alpha1.0-penaltyTrue-l0.5

Overview Model metrics System metrics Traces Artifacts

Metric	Value
val_accuracy	0.7440366972477064
val_macro_precision	0.7399955492294332
val_macro_recall	0.737923337510556
val_macro_f1	0.7382106520182593
val_weighted_precision	0.7503255331944306
val_weighted_recall	0.7440366972477064
val_weighted_f1	0.7463714841278819
final_train_loss	0.6040173675768644
train_loss_curve	0.604063111288051
test_accuracy	0.7300073367571533
test_macro_precision	0.7285611156974254
test_macro_recall	0.7248867193936399
test_macro_f1	0.7262173574092321
test_weighted_precision	0.7363585525729361
test_weighted_recall	0.7300073367571533
test_weighted_f1	0.7326984746382088

7. Task 3.2: Register the best model in the MLflow Model Registry

The run chosen on validation is registered as `<student_ID>-a3-model` and moved to **Staging**.

The same thing can be done in the MLflow UI: open the run, click "Register Model", name it `st127031-a3-model`, then in the **Models** tab open the version and change its stage to Staging.


```
In [22]: MODEL_NAME = f"{STUDENT_ID}-a3-model"

try:
    # register the logged model of the best run, using the URI that
    log_model returned
    registered = mlflow.register_model(model_uri=best_run["model_uri"],
name=MODEL_NAME)
    # archive older versions so only the newest one is in Staging
    MlflowClient().transition_model_version_stage(
        name=MODEL_NAME, version=registered.version, stage="Staging",
        archive_existing_versions=True,
    )
    print(f"Registered '{MODEL_NAME}' version {registered.version} and
set stage=Staging")
except Exception as e:
    print(f"Model registration failed: {type(e).__name__}: {e}")
```

```
2026/10/02 17:26:03 INFO mlflow.store.db.utils: Creating initial MLflow
database tables...
2026/10/02 17:26:03 INFO mlflow.store.db.utils: Updating database tables
INFO [alembic.runtime.migration] Context impl SQLiteImpl.
INFO [alembic.runtime.migration] Will assume non-transactional DDL.
Registered 'st127031-a3-model' version 1 and set stage=Staging
```

```
Successfully registered model 'st127031-a3-model'.
Created version '1' of model 'st127031-a3-model'.
C:\Users\ASUS\AppData\Local\Temp\ipykernel_26900\3313994095.py:7:
FutureWarning:
``mlflow.tracking.client.MlflowClient.transition_model_version_stage`` is
deprecated since 2.9.0. Model registry stages will be removed in a future
major release. To learn more about the deprecation of model registry
stages, see our migration guide here:
https://mlflow.org/docs/latest/model-registry.html#migrating-from-stages
MlflowClient().transition_model_version_stage(
```

MLflow UI: model registry. st127031-a3-model version 1 is in **Staging**.



Registered Models > st127031-a3-model >

Version 1

Registered At: 10/02/2026, 04:35:54 PM Last Modified: 10/02/2026, 04:35:54 PM Source Run: batch-alpha1.0-penaltyTrue-10.5

Stage (deprecated): Staging ⓘ

- > Description Edit
- > Tags
- ▼ Schema

Name	Type
Inputs (11)	
Outputs (5)	

I also save the same model into `app/code/model/`, which is the file the Dash app loads:

```
In [23]: os.makedirs("app/code/model", exist_ok=True)
joblib.dump({
    "W": best_model.W,
    "prep": prep,
    "feature_cols": FEATURE_COLS,
    "num_cols": NUM_COLS,
    "two_level_cols": TWO_LEVEL_COLS,
    "multi_level_cols": MULTI_LEVEL_COLS,
    "owner_scale": OWNER_SCALE,
    "common_brands": COMMON_BRANDS,
    "class_labels": CLASS_LABELS,
    "class_price_ranges": CLASS_PRICE_RANGES,
    "k": k,
}, "app/code/model/car_price_classifier.pkl")
print("Saved app/code/model/car_price_classifier.pkl")
```

Saved app/code/model/car_price_classifier.pkl

Quick check: predict raw cars like the web app does

The web app gets a raw brand typed by the user, so it first groups the brand with the saved `common_brands` list (ignoring upper/lower case), then runs `prep.transform(...)`, adds the bias column, and computes `softmax(X @ W)`. I check a few validation rows (so the test set is still used only once) and one made up car with a rare brand typed in lower case.

```
In [24]: bundle = joblib.load("app/code/model/car_price_classifier.pkl")

def predict_raw(raw_df):
    lookup = {b.lower(): b for b in bundle["common_brands"]}
    raw_df = raw_df[bundle["feature_cols"]].copy()
    raw_df["brand"] = raw_df["brand"].apply(lambda b:
lookup.get(str(b).strip().lower(), "Other"))
    X = add_intercept(to_dense(bundle["prep"].transform(raw_df)))
    probs = best_model.softmax(X @ bundle["W"])
    return probs.argmax(axis=1), probs

rows = df_val.iloc[:5]
pred, probs = predict_raw(rows)
check = pd.DataFrame({
    "brand": rows["brand"].values,
    "year": rows["year"].values,
    "selling_price": rows["selling_price"].values,
    "actual_class": y_val[:5],
    "predicted_class": pred,
    "confidence": probs.max(axis=1).round(3),
})
print(check.to_string(index=False))

# a rare brand typed in lower case should be treated as "Other", not
ignored
volvo = pd.DataFrame([{"year": 2016, "km_driven": 60000, "mileage": 15.0,
"engine": 1969.0,
                        "max_power": 190.0, "seats": 5.0, "owner": 0,
"fuel": "Diesel",
                        "transmission": "Automatic", "seller_type":
"Dealer", "brand": "volvo"}])
p_cls, p_prob = predict_raw(volvo)
print(f"\n'volvo' is grouped as: {'Other' if 'Volvo' not in
bundle['common_brands'] else 'Volvo'}")
print(f"Predicted class: {p_cls[0]} ({bundle['class_labels'][p_cls[0]]},
probabilities {p_prob[0].round(3)})")
```

	brand	year	selling_price	actual_class	predicted_class	confidence
	Maruti	2017	370000	1	2	0.561
	Maruti	2016	290000	1	1	0.670
	Maruti	2016	252000	1	1	0.665
	Mahindra	2019	1650000	3	3	0.999
	Tata	2013	250000	0	0	0.576

```
'volvo' is grouped as: Other
Predicted class: 3 (Luxury), probabilities [0.    0.    0.001 0.999]
```

Most of these are correct, but not all of them. That is expected with an accuracy around 0.73. The test report in Section 6 shows where the errors are: the two middle classes are much harder than the two outer ones (F1 0.60 for Mid-range and 0.64 for Premium, vs 0.84 for Budget and 0.82 for Luxury). A middle class has a neighbouring class on both sides, so a car whose price is close to a class boundary is easy to put in the wrong band.

8. Task 3.3: Unit tests, CI/CD and deployment

This part lives in the repository, not in the notebook:

- **Unit tests** (`app/tests/test_model.py` , run with `pytest`):
 1. `test_model_accepts_expected_input` : the model accepts a DataFrame with the expected columns, also with missing values, an unknown brand and a lowercase brand, and raises an error when a required column is missing.
 2. `test_model_output_shape` : the predicted classes have shape `(m,)` with values 0 to 3, and the probabilities have shape `(m, 4)` with each row adding up to 1.
- **GitHub Actions** (`.github/workflows/ci-cd.yml`): every push runs the unit tests. On a push to `main` , and only if the tests pass, a second job builds the Docker image of the app and pushes it to Docker Hub (`naweeep/car-price-classifier`).
- **Deployment on the ml-brain server.** GitHub cannot SSH into `ml-brain` because it is only reachable from inside CSIM, so the server pulls instead. `app/docker-compose.yaml` runs the app and a small `updater` container that checks Docker Hub every 5 minutes and restarts the app when there is a new image. A commit that fails the tests never gets a new image, so it never reaches the live site.

```
git push
```

```
-> GitHub Actions: pytest
-> tests pass: build the image, push it to Docker Hub
-> ml-brain updater (every 5 min): pull and restart
```

The server no longer runs Traefik (the `web` network from A2 is gone), so the app is published on its own port: <http://ml.brain.cs.ait.ac.th:8031> (inside CSIM). The web app keeps the A1 and A2 pages and adds the A3 classifier page, which shows the predicted price class and the probability of each class.

9. Summary

- **Data leakage:** duplicates were removed before splitting (0 test rows also in train). Brand grouping, class edges, imputers and scaler were fitted on train only. All model comparisons and the MLflow model selection used the validation set, and the test set was used once, on the chosen model (Section 6).
- **Task 1:** `selling_price` was put into 4 classes of similar size (Budget, Mid-range, Premium, Luxury) with `pd.qcut`. The metrics are methods of the `LogisticRegression` class and give the same numbers as `sklearn.metrics.classification_report`.
- **Task 2:** the Ridge penalty can be turned on with `use_penalty=True` and it makes the weights smaller. A small or medium lambda (0.01 to 1) gives about the same validation score as no penalty. The best run used lambda = 0.5 (macro F1 0.738 vs 0.736 without penalty), but this gap is very small, so I would not say ridge clearly helps here. A strong penalty does hurt: lambda = 5 and 10 drop the macro F1 to 0.723 and 0.717, because the weights are pushed too close to zero (underfitting).
- **Task 3:** 9 runs were logged to MLflow (`st127031-a3`). The best run by validation `macro_f1` was registered as `st127031-a3-model` in Staging. Final test result: accuracy 0.730, macro F1 0.726, close to the validation score. The Dash app uses `W`, `prep` and the saved brand list from this model. GitHub Actions runs the unit tests on every push, and only when they pass it pushes a new image that the server picks up by itself (Section 8).